

# EE194 Final Project Ideas

Riley Wood

March 30, 2016

## Possible Research Questions

### 1. Efficiency of Legacy and Modern Concurrent Programming Languages

- With multiprocessor computers becoming the norm, it has been necessary for programming languages to adopt concurrency constructs. Languages like C/C++ generally use locks to protect against data races and deadlock. However, locks in these languages suffer from needing to be globally scoped; abstracting away a lock in a library function means the programmer risks deadlocking the program. New languages like Rust address these shortcomings with different constructs. Specifically, Rust introduces the concept of ownership of data: data can only be owned and thus mutated in one place at a time, thereby avoiding data races and deadlock, but also allowing for abstraction. The question is, do these new constructs slow down execution at all? I would like to run concurrent benchmarks implemented in C and Rust through SNIPER and look for any significant differences in performance, in an effort to discover what if any differences in performance arise because of Rust's new concurrency constructs.

### 2. When To Switch From Ring Cache to Mesh Network Cache?

- Multicore systems often organize their L2 caches in a ring structure. This means coherence messages only need to hop at most  $\frac{N}{2}$  times to reach their destination, where  $N$  is the number of cores. At a certain point, though, the number of cores grows the cache coherence overhead too much, and it is preferable to use a mesh network. I would be interested in seeing at what point a mesh network becomes the preferable architecture: how many CPUs under which workloads with what other combinations of system parameters make a ring no longer viable?

### 3. Comparing Cache Coherence Protocols

- I would like to see how various cache coherence protocols perform relative to one another across several workloads. I would look at protocols like MESI, MSI, MOSI, MOESI, MERSI, MESIF, write-once, and others. I would compare them based on metrics such as IPC, average memory miss penalty, average memory access time, misses per 1K instructions, and other memory-focused metrics.